

Predicting Billboard #1 Hits

Authors: May Eindra Tet Toe, Anastasia Tountas, Tran Anh Thu Phung, Harry Nguyen

Summary

Insert a summary of your analysis here.

Introduction

Insert an introduction to your analysis here.

Methods and Results

Overview of our Methodology

The Billboard Hot 100 dataset provides us with two datasets: `billboard.csv`, which contains one row per song that reached #1 along with its features, and `topics.csv`, a reference list of lyrical theme labels (`lyrical_topics`). Our team decided to combine the two datasets to `billboard_clean` in order to create more comprehensive data to better equip our model. Here is an overview of our data analysis methodology:

1. **Reading & Wrangling:** read and wrangle the dataset into one tidy combined dataset
2. **Exploratory Data Analysis:** explore the data to understand the distribution of features and their relationships with the target variable
3. **Model Building:** build a regression model to predict the number of weeks a song stays at #1 based on its features
4. **Model Evaluation:** evaluate the performance of our model using appropriate metrics and validation techniques
5. **Results and Conclusion:** summarize our findings, discuss the implications of our results, and suggest potential future work or improvements to our analysis

Load libraries and set seed

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.4.3
```

```
## Warning: package 'ggplot2' was built under R version 4.4.3
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
```

```
## v forcats    1.0.0      v stringr    1.5.1
```

```
## v ggplot2    4.0.2      v tibble     3.2.1
```

```
## v lubridate  1.9.3      v tidyr      1.3.1
```

```
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(tidyuesdayR)
```

```
## Warning: package 'tidytuesdayR' was built under R version 4.4.3
```

```
library(knitr)
```

```
library(dplyr)
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```
##
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

```
library(tidymodels)
```

```
## Warning: package 'tidymodels' was built under R version 4.4.3
```

```
## -- Attaching packages ----- tidymodels 1.4.1 --
```

```
## v broom      1.0.12      v rsample      1.3.2
```

```
## v dials      1.4.2       v tailor       0.1.0
```

```
## v infer      1.1.0       v tune         2.0.1
```

```
## v modeldata  1.5.1       v workflows    1.3.0
```

```
## v parsnip    1.4.1       v workflowsets 1.1.1
```

```
## v recipes    1.3.1       v yardstick    1.3.2
```

```
## Warning: package 'broom' was built under R version 4.4.3
```

```
## Warning: package 'dials' was built under R version 4.4.3
```

```
## Warning: package 'scales' was built under R version 4.4.3
```

```
## Warning: package 'infer' was built under R version 4.4.3
```

```
## Warning: package 'modeldata' was built under R version 4.4.3
```

```
## Warning: package 'parsnip' was built under R version 4.4.3
```

```
## Warning: package 'recipes' was built under R version 4.4.3
```

```
## Warning: package 'rsample' was built under R version 4.4.3
```

```
## Warning: package 'tailor' was built under R version 4.4.3
```

```
## Warning: package 'tune' was built under R version 4.4.3
```

```
## Warning: package 'workflows' was built under R version 4.4.3
```

```
## Warning: package 'workflowsets' was built under R version 4.4.3
```

```
## Warning: package 'yardstick' was built under R version 4.4.3
```

```
## -- Conflicts ----- tidymodels_conflicts() --
```

```
## x rsample::calibration() masks caret::calibration()
```

```
## x scales::discard() masks purrr::discard()
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x recipes::fixed() masks stringr::fixed()
```

```
## x dplyr::lag() masks stats::lag()
```

```
## x caret::lift() masks purrr::lift()
```

```
## x yardstick::precision() masks caret::precision()
```

```
## x yardstick::recall() masks caret::recall()
```

```
## x yardstick::sensitivity() masks caret::sensitivity()
## x yardstick::spec() masks readr::spec()
## x yardstick::specificity() masks caret::specificity()
## x recipes::step() masks stats::step()

set.seed(123)
```

Reading & Wrangling our dataset from the web into R

- The two datasets were joined using a `left_join()` on `lyrical_topic` (billboard) and `lyrical_topics` (topics).
- **Identifier columns** (`song`, `artist`, `date`) were removed as they are not predictive features.
- **non_consecutive** was removed to prevent data leakage — it can only be determined after a song has completed its full chart run.
- **rating_1**, **rating_2**, **rating_3** were removed in favour of **overall_rating** (their mean) to avoid multicollinearity.
- **Free-text columns** were removed to avoid the high dimensionality and noise associated with Natural Language Processing: `lyrics`, `u_s_artwork`, `sound_effects`, `song_structure`, `featured_artists`, `songwriters`, `songwriters_w_o_interpolation_sample_credits`, `producers`, `double_a_side`, and `talent_contestant`.
- **High-cardinality columns** removed because they generalise poorly: `label`, `parent_label`, and `artist_place_of_origin`.
- **Redundant columns** were removed: `keys` is replaced by `simplified_key`; `cdr_style` and `discogs_style` are replaced by `cdr_genre` and `discogs_genre` respectively.
- **discogs_genre** was removed in favour of `cdr_genre` because `cdr` has more reliable class definitions, thus improving internal validity.
- **Ambiguous columns** were removed: `featured_in_a_then_contemporary_play`, `featured_in_a_then_contemporary_t_v_show`.
- All remaining character columns were converted to factors using `as.factor()`, and rows with any missing values were dropped with `drop_na()`.

```
# Load the dataset
tuesdata <- tidyTuesdayR::tt_load(2025, week = 34)

## ---- Compiling #TidyTuesday Information for 2025-08-26 ----
## --- There are 2 files available ---
##
##
## -- Downloading files -----
##
## 1 of 2: "billboard.csv"
## 2 of 2: "topics.csv"

# Extract the datasets
billboard <- tuesdata$billboard
topics <- tuesdata$topics

# Join the two datasets
billboard_joined <- billboard |>
  left_join(topics, by = c("lyrical_topic" = "lyrical_topics"))

billboard_clean <- billboard_joined |>
  # Remove identifier columns
  select(-c(song, artist, date)) |>
  # Remove data-leakage column
  select(-non_consecutive) |>
  # Remove individual judge scores to avoid multicollinearity
```

```

select(-c(rating_1, rating_2, rating_3)) |>
# Remove free-text columns
select(-c(
  lyrics,
  u_s_artwork,
  sound_effects,
  song_structure,
  featured_artists,
  songwriters,
  songwriters_w_o_interpolation_sample_credits,
  producers,
  double_a_side,
  talent_contestant
)) |>
# Remove high-cardinality columns
select(-c(label, parent_label, artist_place_of_origin)) |>
# Remove redundant columns and discogs_genre
select(-c(cdr_style, discogs_style, keys, discogs_genre)) |>
# Remove ambiguous columns
select(-c(
  featured_in_a_then_contemporary_play,
  featured_in_a_then_contemporary_film,
  featured_in_a_then_contemporary_t_v_show
)) |>
# Convert remaining character columns to factors
mutate(across(where(is.character), as.factor)) |>
drop_na()

```

Exploratory Data Analysis

First, we find the summary statistics for our target variable, `weeks_at_number_one`, and our key numeric predictors: `overall_rating`, `bpm`, `energy`, `danceability`, `happiness`, `acousticness`, `loudness_d_b`, `front_person_age`, and `length_sec`.

```

## Warning in attr(x, "align"): 'xfun::attr()' is deprecated.
## Use 'xfun::attr2()' instead.
## See help("Deprecated")

## Warning in attr(x, "format"): 'xfun::attr()' is deprecated.
## Use 'xfun::attr2()' instead.
## See help("Deprecated")

```

Table 1: Summary statistics

Variable	Mean	SD	Min	Median	Max
acousticness	31.51	27.66	0	23.00	98
bpm	115.50	23.69	16	115.00	176
danceability	62.45	15.21	14	64.00	98
energy	59.81	19.62	3	60.50	98
front_person_age	27.97	6.51	11	27.50	62
happiness	62.60	24.96	4	68.00	99
length_sec	222.88	50.12	96	224.00	515
loudness_d_b	-9.02	3.40	-23	-9.00	-1
overall_rating	6.18	1.86	1	6.33	10

Variable	Mean	SD	Min	Median	Max
weeks_at_number_one	2.91	2.49	1	2.00	16

Then, we examine the distribution of our target variable, `weeks_at_number_one`, to check for skewness.




Figure 1: Distribution of weeks at number one.

The distribution is heavily right-skewed. The majority of songs spent only 1 to 3 weeks at #1, with counts dropping off sharply afterwards. A small number of songs stayed for 10 or more weeks, forming a long right tail.

Next, we examine the distributions of the five Spotify audio features, to identify any skewness that may require transformation before modelling.

`acousticness` is heavily right-skewed, with most songs clustered near zero. `happiness` shows a bimodal pattern with peaks around 30 and around 80. `energy` is very slightly left-skewed, while `bpm` and `danceability` are roughly normal.

Finally, we compare average `weeks_at_number_one` by genre to see whether certain genres are associated with longer chart runs.

Electronic/Dance;Latin seems to have the highest average (~10 weeks). Pop;Reggae and Folk/Country;March also rank highly. Mainstream genres like Pop and Rock sit in the middle of the range. At the lower end, we can find niche genres like Electronic/Dance;Funk/Soul, averaging to 1 week.




Figure 2: Distributions of the five Spotify audio features.




Figure 3: Average weeks at number one by genre.

Feature Transformations

Based on the Exploratory Data Analysis, two variables require transformation before modelling:

- **weeks_at_number_one** is heavily right-skewed
- **acousticness** is also heavily right-skewed

Therefore, we apply log transformation to these three variables to compresses the tail and produce a more symmetric target distribution. We will be using `log1p()` instead of `log()` to handle any zero values in the data.

```
billboard_log_transformed <- billboard_clean |>
  mutate(
    log_weeks_at_number_one = log1p(weeks_at_number_one),
    log_acousticness = log1p(acousticness)
  )
```

Moreover, several songs in the dataset were associated with more than one musical genre (for example, Electronic/Dance; Latin or Pop; Funk/Soul). Treating these combined labels as single categorical values would incorrectly imply that each combination represents a distinct genre and would lead to sparse, poorly interpretable factor levels. To address this, the `cdr_genre` variable was processed as a multi-label feature. Individual genre were first separated into their component parts and then encoded as binary indicator variables using one-hot encoding.

```
billboard_genres <- billboard_log_transformed |>
  separate_rows(cdr_genre, sep = ";") |>
  mutate (cdr_genre = str_trim(cdr_genre))

billboard_genres_wide <- billboard_genres |>
  mutate(value = 1) |>
  pivot_wider(
    names_from = cdr_genre,
    values_from = value,
    values_fill = 0
  )
```

Transform the Simplified Key:

```
billboard_simplify_key <- billboard_genres_wide |>
  mutate(
    simplified_key = case_when(
      simplified_key == "Multiple Keys" ~ "Multiple Keys",
      str_ends(simplified_key, "m") ~ "Minor", # Am, Bbm, Bm, Cm, etc.
      TRUE ~ "Major" # A, Ab, B, Bb, C, etc.
    ) |> as.factor()
  )

billboard_model_final <- billboard_simplify_key |>
  select(-weeks_at_number_one, -acousticness)
```

Feature Selection

Drop Redundant Columns

- Demographic identity columns: Race and gender of the artist or producing crews don't directly cause a song to stay longer at #1, so we should drop these columns.
- Songwriter/producer role overlaps
- Timing collinearity

- Lyric redundant content
- External content

```
columns_to_drop <- c(
  # Demographic identity
  "artist_male", "artist_white", "artist_black", "songwriter_male", "songwriter_white", "producer_male"
  # Artist structure
  "artist_structure", "group_named_after_non_lead_singer",
  # Songwriter/producer overlaps
  "artist_is_only_songwriter", "artist_is_only_producer", "songwriter_is_a_producer",
  # Timing
  "instrumental_length_sec", "intro_length_sec",
  # Lyrical content
  "lyrical_topic", "lyrical_narrative",
  # External content
  "written_for_a_play", "written_for_a_film", "written_for_a_t_v_show", "eurovision_entry", "topped_the

billboard_manual_dropped <- billboard_model_final |>
  select(-all_of(columns_to_drop))
```

```
# Identify and remove near-zero variance predictors
nzv_cols <- nearZeroVar(billboard_manual_dropped,
  freqCut = 95/5, # flag if most common value appears 19x more than second most
  uniqueCut = 5, # flag if fewer than 5% unique values
  names = TRUE)

nzv_cols
```

Drop columns with nearly zero variance :

```
## [1] "posthumous"          "time_signature"
## [3] "accordion"           "banjo"
## [5] "bongos"              "clarinet"
## [7] "cowbell"             "flute_piccolo"
## [9] "harmonica"           "human_whistling"
## [11] "kazoo"               "mandolin"
## [13] "pedal_lap_steel"     "ocarina"
## [15] "sitar"               "trumpet"
## [17] "ukulele"             "violin"
## [19] "rap_verse_in_a_non_rap_song" "instrumental"
## [21] "free_time_vocal_introduction" "live"
## [23] "foreign_language"    "Folk/Country"
## [25] "March"              "Jazz"
## [27] "Polka"              "Reggae"
## [29] "Calypso"            "Blues"
## [31] "Latin"
```

```
billboard_model_selected <- billboard_manual_dropped |>
  select(-all_of(nzv_cols))
```

Regression Model

```
# Train/test split
billboard_split <- initial_split(billboard_model_selected, prop = 0.7)
train_data <- training(billboard_split)
```

```

test_data <- testing(billboard_split)

lm_recipe <- recipe(log_weeks_at_number_one ~ ., data = train_data) |>

  # Convert any remaining factor predictors to dummy variables
  # (simplified_key gets 3 dummies; reference level = "Major")
  step_dummy(all_nominal_predictors(), one_hot = FALSE) |>

  # Remove any predictor with near-zero variance that survived earlier filtering
  step_nzv(all_predictors()) |>

  # Centre and scale all numeric predictors for coefficient interpretability
  step_normalize(all_numeric_predictors()) |>

  # Drop any predictor that is a perfect linear combination of others
  step_lincomb(all_numeric_predictors())

# Inspect the processed training frame (sanity check)
prep(lm_recipe) |> bake(new_data = NULL) |> glimpse()

## Rows: 686
## Columns: 38
## $ overall_rating      <dbl> 0.06012931, -0.11999606, -1.02062290, 0.7~
## $ divisiveness        <dbl> 0.4989845, -1.5882310, 0.4989845, -0.1967~
## $ multiple_lead_vocalists <dbl> -0.5286235, -0.5286235, -0.5286235, 1.888~
## $ front_person_age     <dbl> -0.29475019, 1.08748951, -1.36982551, 0.4~
## $ artist_is_a_songwriter <dbl> 0.763835, 0.763835, -1.307275, 0.763835, ~
## $ artist_is_a_producer  <dbl> -0.725996, -0.725996, -0.725996, -0.72599~
## $ bpm                 <dbl> -0.1899626, -0.3157082, -0.1899626, -1.78~
## $ energy              <dbl> 0.30537174, 1.29670806, 0.04449377, -0.79~
## $ danceability         <dbl> 0.71918960, 0.51858074, -0.81881169, -0.6~
## $ happiness            <dbl> 0.77203245, 0.73140851, -1.29978831, -1.3~
## $ loudness_d_b         <dbl> 0.5937448, 0.8908337, 0.2966559, -1.78296~
## $ vocally_based        <dbl> -0.2419845, -0.2419845, -0.2419845, -0.24~
## $ bass_based           <dbl> 2.018795, -0.494623, -0.494623, -0.494623~
## $ guitar_based         <dbl> 0.822359, 0.822359, 0.822359, -1.214241, ~
## $ piano_keyboard_based  <dbl> 0.9453692, -1.0562458, 0.9453692, 0.94536~
## $ orchestral_strings    <dbl> -0.6660052, -0.6660052, 1.4993009, 1.4993~
## $ horns_winds          <dbl> -0.6682854, 1.4941853, -0.6682854, -0.668~
## $ falsetto_vocal       <dbl> -0.4461056, -0.4461056, -0.4461056, -0.44~
## $ handclaps_snaps      <dbl> -0.3907918, 2.5551771, -0.3907918, -0.390~
## $ saxophone            <dbl> -0.2519389, 3.9634297, -0.2519389, -0.251~
## $ length_sec           <dbl> -0.16603410, 3.23164268, -1.47877286, 3.0~
## $ vocal_introduction    <dbl> -0.2771479, -0.2771479, -0.2771479, -0.27~
## $ fade_out             <dbl> -1.2441568, 0.8025856, 0.8025856, -1.2441~
## $ cover                <dbl> -0.4176158, -0.4176158, -0.4176158, -0.41~
## $ sample               <dbl> -0.2740902, -0.2740902, -0.2740902, -0.27~
## $ interpolation         <dbl> -0.3341683, -0.3341683, -0.3341683, -0.33~
## $ inspired_by_a_different_song <dbl> -0.647825, -0.647825, -0.647825, -0.64782~
## $ spoken_word          <dbl> -0.3036808, -0.3036808, -0.3036808, -0.30~
## $ explicit             <dbl> -0.4484501, -0.4484501, -0.4484501, -0.44~
## $ log_acousticness      <dbl> -1.18809397, 0.66158181, -0.25132153, 0.0~
## $ Pop                  <dbl> -0.725996, -0.725996, -0.725996, 1.375410~
## $ Rock                 <dbl> 1.5913471, -0.6274824, 1.5913471, -0.6274~

```

```
## $ `Funk/Soul`          <dbl> -0.5308781, 1.8809257, -0.5308781, -0.530~
## $ `Electronic/Dance`  <dbl> -0.2647649, 3.7714294, -0.2647649, -0.264~
## $ `Hip Hop`           <dbl> -0.2950189, -0.2950189, -0.2950189, -0.29~
## $ log_weeks_at_number_one <dbl> 0.6931472, 1.0986123, 1.0986123, 1.609437~
## $ simplified_key_Minor <dbl> -0.584779, 1.707555, -0.584779, -0.584779~
## $ simplified_key_Multiple.Keys <dbl> -0.3177675, -0.3177675, -0.3177675, -0.31~
```

```
lm_spec <- linear_reg() |>
  set_engine("lm") |>
  set_mode("regression")

lm_workflow <- workflow() |>
  add_recipe(lm_recipe) |>
  add_model(lm_spec)
```

```
set.seed(123)
cv_folds <- vfold_cv(
  train_data,
  v      = 5,
  repeats = 3,
  strata = log_weeks_at_number_one
)

cv_results <- fit_resamples(
  lm_workflow,
  resamples = cv_folds,
  metrics    = metric_set(rmse, rsq, mae),
  control    = control_resamples(save_pred = TRUE)
)

# Cross-validation summary
cv_metrics <- collect_metrics(cv_results)
print(cv_metrics)
```

```
## # A tibble: 3 x 6
##   .metric .estimator  mean     n std_err .config
##   <chr>   <chr>      <dbl> <int>   <dbl> <chr>
## 1 mae     standard    0.415     15 0.00599 pre0_mod0_post0
## 2 rmse     standard    0.505     15 0.00674 pre0_mod0_post0
## 3 rsq      standard    0.0626    15 0.0113  pre0_mod0_post0
```

```
final_fit <- fit(lm_workflow, data = train_data)
```

```
test_predictions <- predict(final_fit, new_data = test_data) |>
  bind_cols(test_data |> select(log_weeks_at_number_one)) |>
  rename(
    predicted = .pred,
    actual    = log_weeks_at_number_one
  )

test_metrics <- test_predictions |>
  metrics(truth = actual, estimate = predicted)

print(test_metrics)
```

```
## # A tibble: 3 x 3
```

```
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 rmse    standard    0.502
## 2 rsq     standard    0.0518
## 3 mae     standard    0.395
```

Result Visualization

```
# Pre-compute R2 as a plain string to avoid encoding issues with the superscript
rsq_label <- paste0(
  "R_squared = ",
  round(
    test_predictions |>
      metrics(truth = actual, estimate = predicted) |>
      filter(.metric == "rsq") |>
      pull(.estimate),
    3
  )
)

p_actual_vs_predicted <- test_predictions |>
  mutate(
    actual_weeks = expm1(actual),
    predicted_weeks = expm1(predicted)
  ) |>
  ggplot(aes(x = actual_weeks, y = predicted_weeks)) +
  geom_point(alpha = 0.4, colour = "steelblue", size = 2) +
  geom_abline(
    slope = 1,
    intercept = 0,
    linetype = "dashed",
    colour = "firebrick",
    linewidth = 0.8
  ) +
  annotate(
    "text",
    x = 12,
    y = 2,
    label = rsq_label, # plain ASCII string - no special characters
    colour = "firebrick",
    size = 4
  ) +
  labs(
    title = "Actual vs Predicted Weeks at #1 (Test Set)",
    subtitle = "Dashed line = perfect prediction; points below = underestimate",
    x = "Actual Weeks at #1",
    y = "Predicted Weeks at #1"
  ) +
  theme_minimal()

print(p_actual_vs_predicted)
```

```
# Extract coefficients safely using broom::tidy explicitly
coef_data <- broom::tidy(extract_fit_engine(final_fit)) |>
```




Figure 4: Actual vs predicted weeks at number one on the test set (original scale).

```

filter(term != "(Intercept)") |>
arrange(desc(abs(estimate))) |>
slice_head(n = 15) |>
mutate(
  direction = if_else(estimate > 0, "Positive", "Negative"),
  term       = str_replace_all(term, "_", " ") |> str_to_title()
)

p_coefficients <- ggplot(
  coef_data,
  aes(x = reorder(term, abs(estimate)), y = estimate, fill = direction)
) +
  geom_col(show.legend = TRUE) +
  coord_flip() +
  scale_fill_manual(
    values = c("Positive" = "steelblue", "Negative" = "firebrick"),
    name   = "Effect on\nlog(weeks + 1)"
  ) +
  labs(
    title    = "Top 15 Predictors by Coefficient Magnitude",
    subtitle = "Standardised coefficients - larger absolute value = stronger effect",
    x        = NULL,
    y        = "Standardised Coefficient"
  ) +
  theme_minimal() +
  theme(legend.position = "right")

print(p_coefficients)

```




Figure 5: Top 15 predictors by standardised coefficient magnitude.